

MODEL CHECKING USING SPIN/PROMELA

courtsey:

A K Bhattacharjee, BARC, Mumbai; CFDVS, IIT, Bombay

Scott Smolka, Stony Brook, USA

Theo C Ruys, Univ of Twente

Gerard Holzmann, Design and Validation of Computer Protocols

Formal Verification : Background

- The Process of Formal Verification consists of

- Requirements Capture
 - Modeling
 - Specification
 - Analysis
- Documentation



- In reality, some phases may overlap

- the overall process is iterative rather than sequential

Model Checking

- Model checking
 - is one of the powerful FORMAL VERIFICATION technique.
 - allows one to verify temporal properties of a finite state representation.
- The finite state representation is that of a typical concurrent system.....
- What do we mean by representation here ?
- The representation is a model of the system.....

Model Checking.....

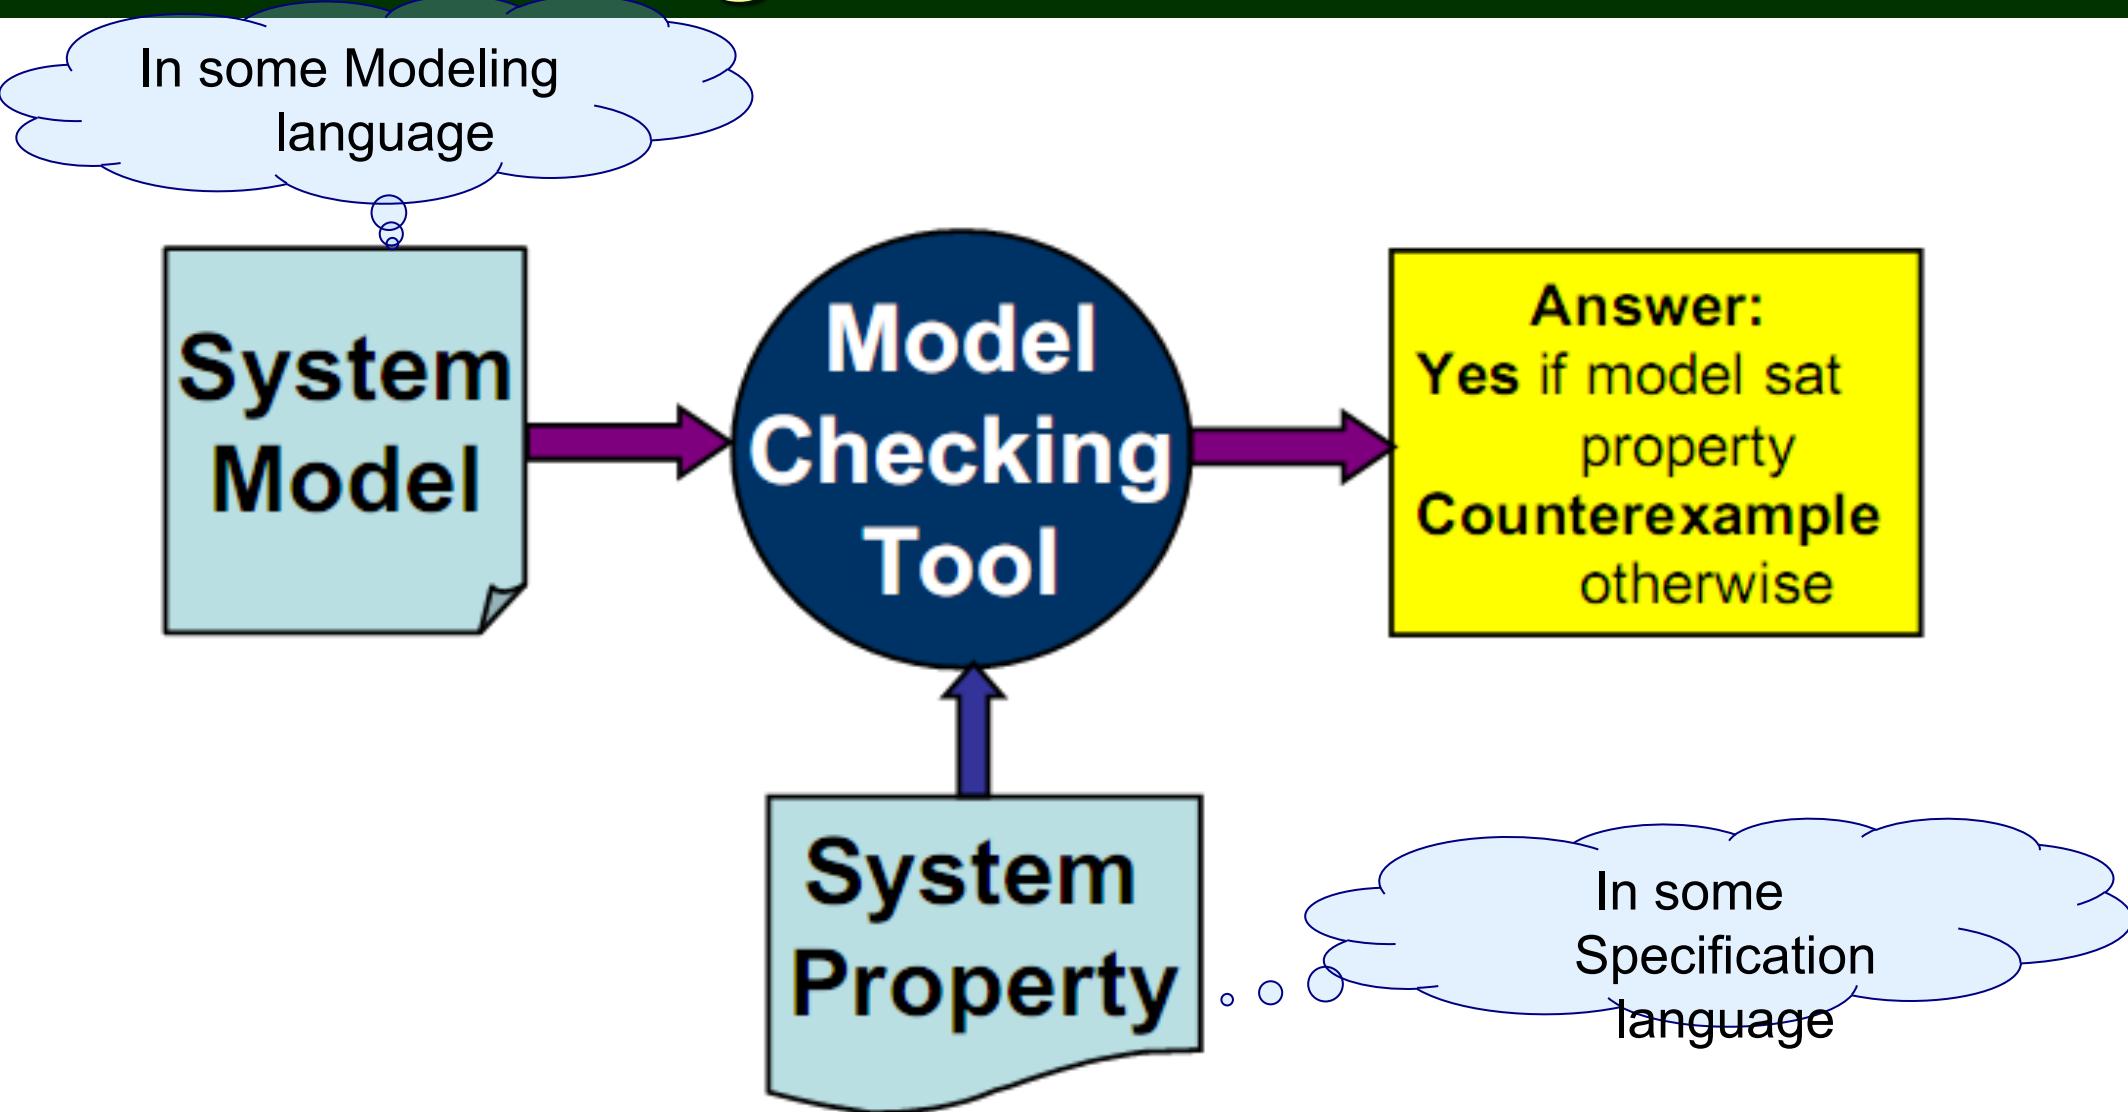
- The basic idea is
 - that a **finite state model** of a system is **systematically explored** in order to determine whether or not a given **temporal property holds**.
 - deliver **a counter example** if the specified property does not hold.
- State space explosion
 - because a complete state space is generated for a given model, the **analysis may fail due lack of memory** if the model is too large
 - can be tackled **via abstraction**.

Model Checking

- **Verifying design models**
 - against specification for finite state concurrent systems.
- **It is an automated technique wherein**
 - Inputs
 - finite state model of the system & properties stated in some standard formalism
 - Outputs
 - property valid against the model or not.
- **Many Commercial and academic tools**
 - Spin (Bell Labs.), Formal-Check (Cadence), VIS (UCB), SMV (CMU, Cadence)
 - In-house tools: Rule Base (IBM), Intel, SUN, Bingo (Fujitsu), etc

Model Checking

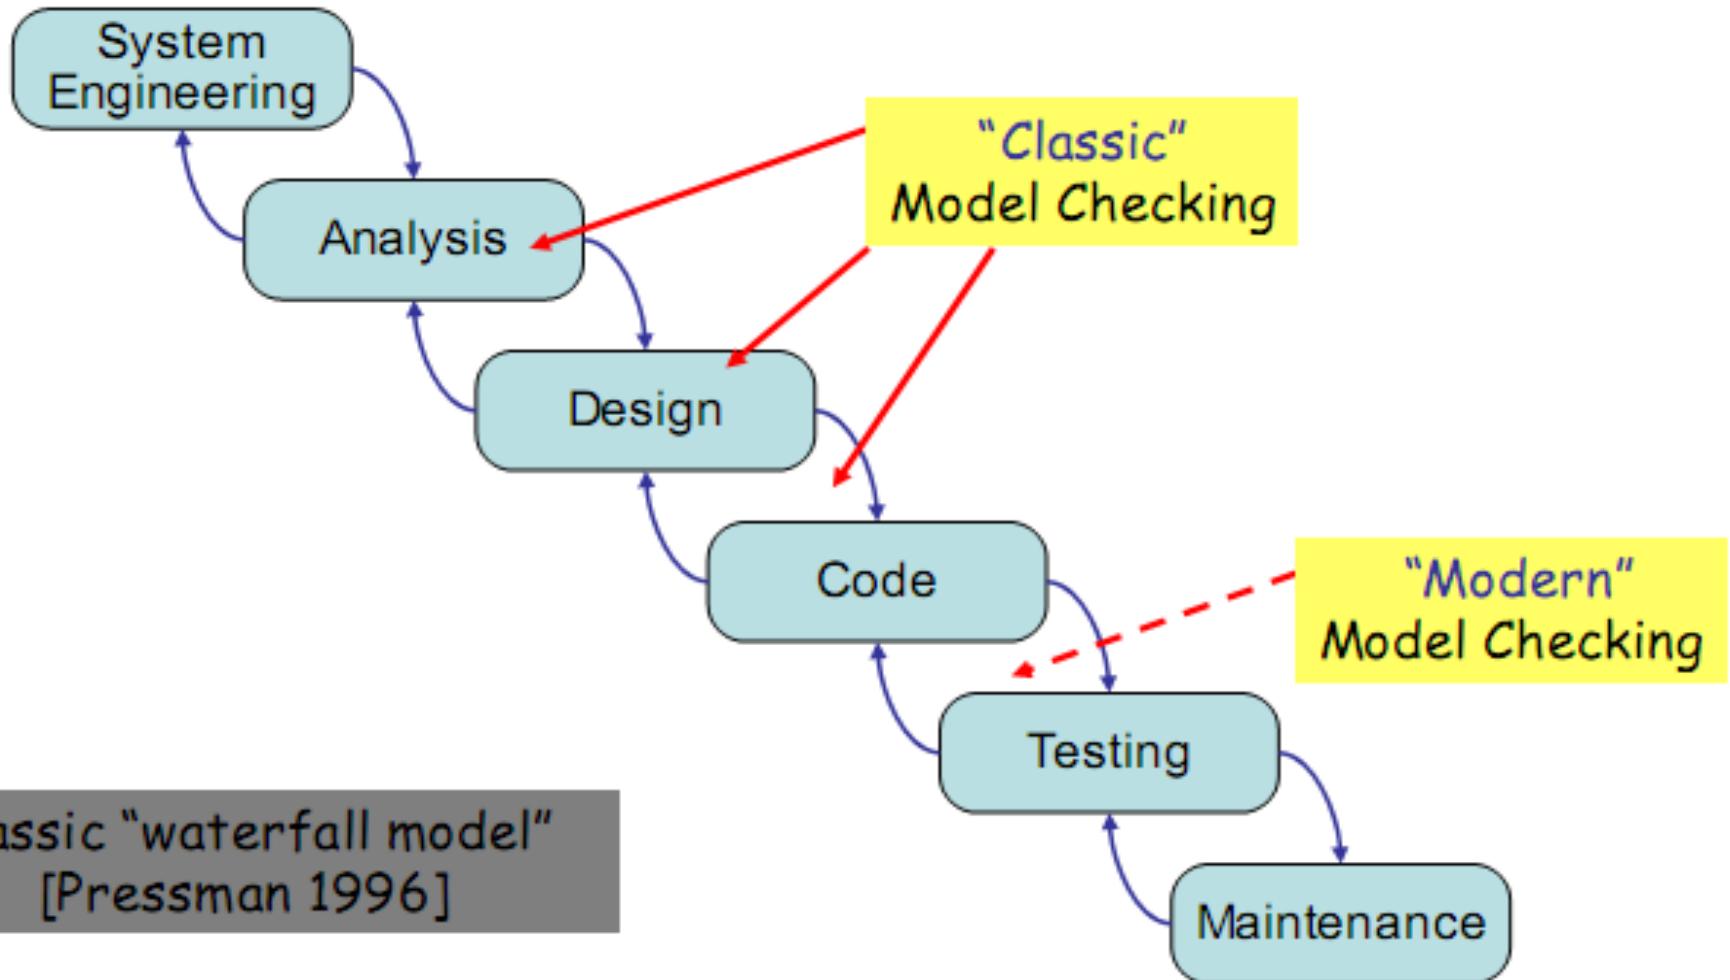
[Smolka-StonyBrook]



The entire process is automated

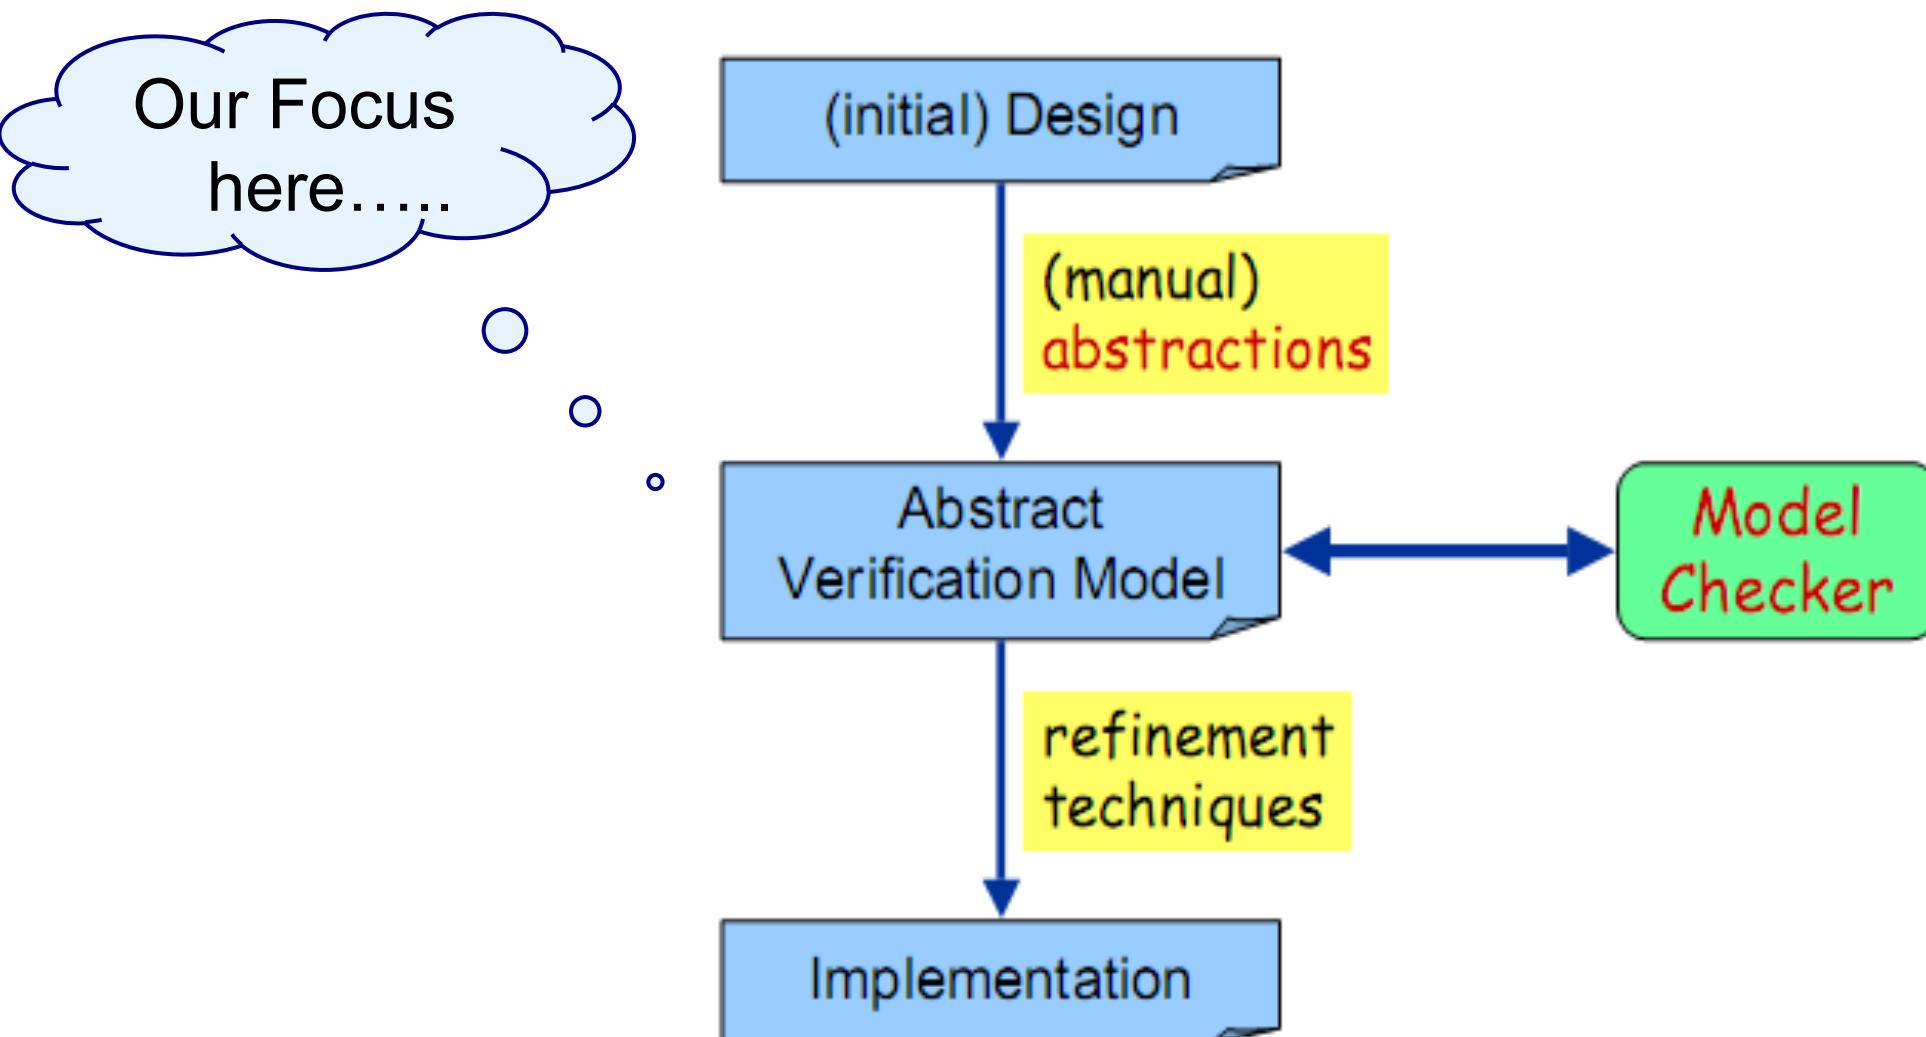
System Development

[Smolka-StonyBrook]



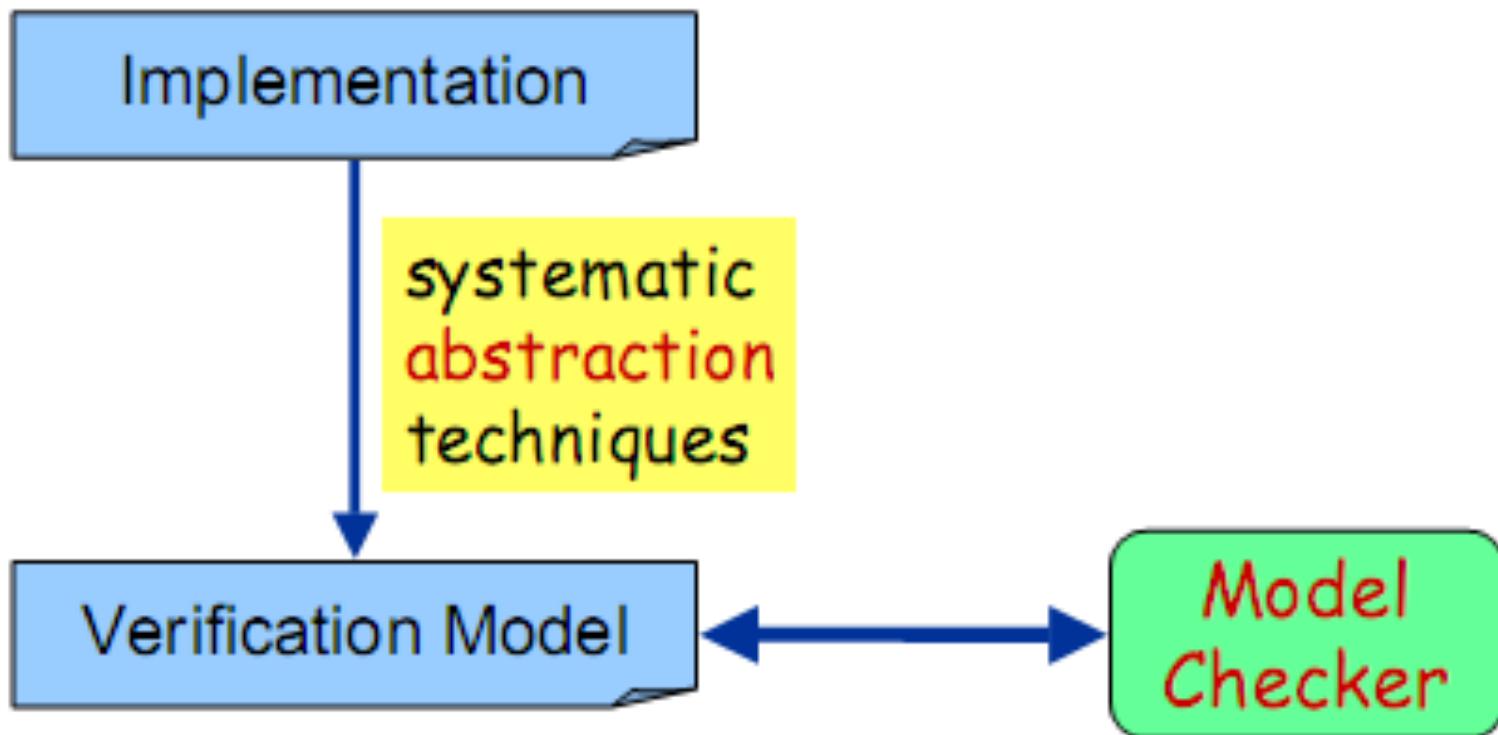
Classic Model Checking

[Theo Ruys]



Modern Model Checking

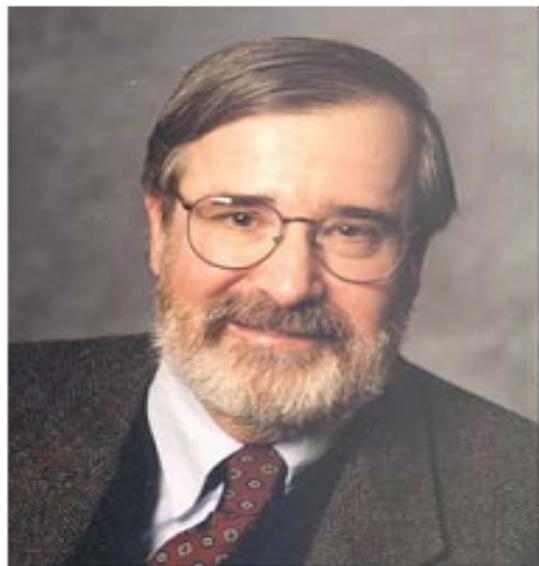
[Theo Ruys]



Who developed it?

- Developed independently by Clarke, Emerson, and Sistla and by Queille and Sifakis in early 1980's.

Clarke Emerson Sifakis Receive 2007 Turing Award



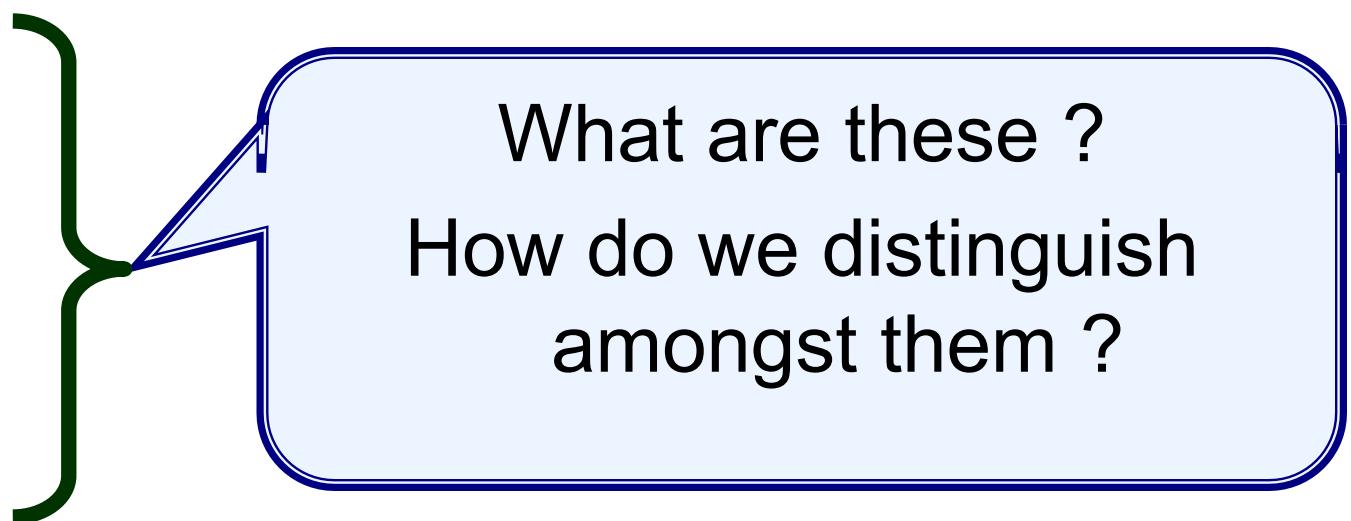
... they developed this fully automated approach [Model Checking] that is now the most widely used verification method in the hardware and software industries.

Applicability : Distributed Systems

- Specific concern on the distributed systems
 - Network applications
 - Data communication protocols
 - Multithreaded code
 - Client-Server applications
- Suffer from common design flaws

Common Design Flaws

- Deadlock
- Livelock,
- Starvation
- Race conditions...



Common Design Flaws.....

- Deadlock
- Livelock, Starvation
- Underspecification
 - unexpected reception of messages
- Overspecification
 - Dead code
- Violations of constraints
 - Buffer overruns
 - Array bounds violations
- Assumptions about speed
 - Logical correctness vs. real-time performance

Model Checking Definition

[Clarke & Emerson 1981]

- “Model checking is an automated technique that,
 - given a finite-state model of a system and
 - a logical property, systematically checks whether this property holds for (“a given initial state in) that model.”

$$M \models \varphi ?$$

Does system model **M** satisfy system property **φ**?

M given as a **state machine**, that is Finite-state
φ usually specified in **temporal logic**.

Model Checking with SPIN

- Involves Three Steps
- modeling
 - convert the design into a formalism to be accepted by the model checking tool SPIN
- specification
 - state the properties that the design must satisfy
 - must be complete
- verification
 - normally based on a tool i.e on SPIN
 - is automatic
 - analysis of verification results is, however, manual.

SPIN Primary Usage : Verifying Protocols

■ Protocol definition

- a service specification
- explicit assumptions about the environment
- the protocol vocabulary
- formal definition
- procedural rules
 - a protocol validation model
 - describes the interaction of processes in the system
 - in some formal language
 - give only a partial description of protocol. Why?

Introduction to SPIN

■ SPIN – Simple Promela INterpreter

- a tool for analysing the **logical consistency** of concurrent systems
 - specifically of data communication protocols
- is a state-of-the-art model checker, used by >2000 users
- system model of a concurrent system is described in PROMELA

■ PROMELA – PROcess MEta Language

- specification language to model **finite-state** systems
- allows for **dynamic creation** of concurrent processes
- expressive enough **to describe processes, their interactions..**
 - in **synchronous** as well as **asynchronous** manner
- resembles the programming language C ...plus CSP...inspired by Dijkstra's guarded commands

Introduction to SPIN...

- Best way to get started with SPIN is via XSPIN.
- XSPIN allows
 - To edit the Promela model
 - Simulate the model (Closed)
 - Random
 - Interactive
 - Guided
- Invoking the verifier
 - exhaustive OR bitstate hashing
- LTL Property Manager
- Help

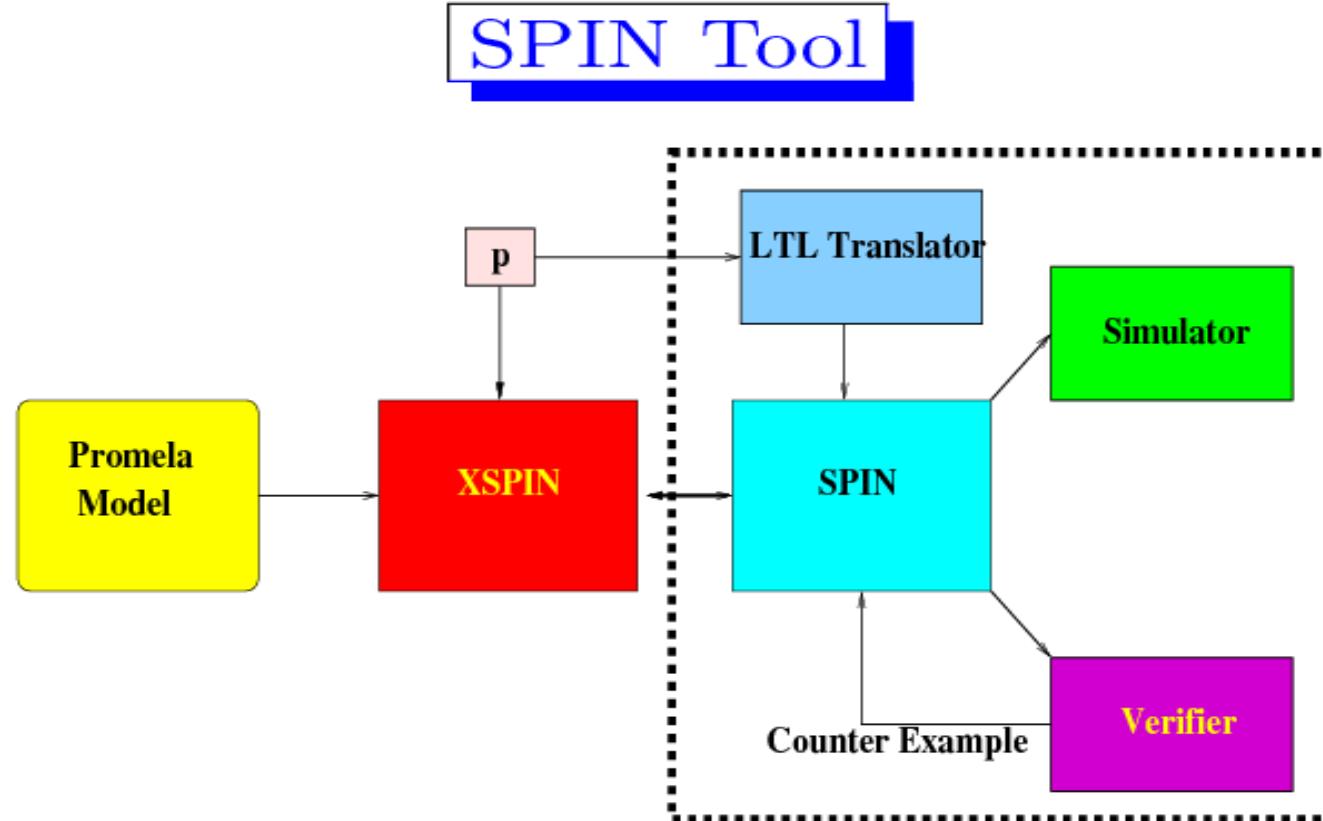
Introduction to SPIN...

- The look-and-feel of XSPIN
 - Main text window
 - basic edit & search capabilities.
 - File handling
 - basic browsing, loading, saving capabilities
 - File editing
 - cut, copy & paste .
 - Tool support
 - syntax checking, simulation, verification, automaton portray.
 - Help facilities

Introduction to SPIN...

- Promela and SPIN/XSPIN are
 - developed by Gerard Holzmann at Bell Labs
 - Freeware for non-commercial use
 - is a State-of-art model checker
 - another is SMV

The SPIN tool



Introduction to SPIN...

[Theo Ruys]

- Major versions:

1.0	Jan 1991	initial version [Holzmann 1991]
2.0	Jan 1995	partial order reduction
3.0	Apr 1997	minimised automaton representation
4.0	late 2002	Ax: automata extraction from C code

- Some success factors of SPIN (subjective!):

- “press on the button” verification (model checker)
- very efficient implementation (using C)
- nice graphical user interface (Xspin)
- not just a research tool, but well supported
- contains more than two decades research on advanced computer aided verification (many optimization algorithms)

Introduction to SPIN...

[Theo Ruys]

- SPIN's starting page:

<http://netlib.bell-labs.com/netlib/spin/whatispin.html>

- Basic SPIN manual
- Getting started with Xspin
- Getting started with SPIN
- Examples and Exercises
- Concise Promela Reference (by Rob Gerth)
- Proceedings of all SPIN Workshops

Also part of SPIN's documentation distribution
(file: `html.tar.gz`)

- Gerard Holzmann's website for papers on SPIN:

<http://cm.bell-labs.com/cm/cs/who/gerard/>

- SPIN version 1.0 is described in [Holzmann 1991].

PROMELA

- defining a validation model consisting of
 - a set of states incorporating info about values of variables, program counters etc
 - a transition relation
- a representation of a FSM in terms of
 - processes
 - message channels
 - state variables
- only the design of consistent set of rules to govern interaction amongst processes in a DS
- NOT implementation details

PROMELA....

- PROMELA model consists of
 - type declarations
 - channel declarations
 - variable declarations
 - process declarations
 - init process
- As mentioned, PROMELA model = FSM (usually a very large)
 - but it is finite, and hence has
 - no unbounded data
 - no unbounded channels
 - no unbounded processes
 - no unbounded process creations.

Processes

- A process type (proctype) consists of
 - a name, list of formal parameters, declarations of local variables body
- A process
 - executes concurrently with all other processes
 - communicates with other processes using channels & global variables
 - may access shared variables
 - defined by proctype declarations
- Each process has
 - its program counter
 - local variables (local state)

There may be
several processes
of the
same type !!!

PROMELA : An illustration

```
proctype hello() {
    printf("Hello\n")
}

proctype world() {
    printf("World\n")
}

init {
    run hello(); run world ();
}
```

Variables and Basic Data Types

■ Promela variables

- provide the means of storing information about the system being modelled.
- may hold global information on the system or information that is local to a particular component (process).
- supports five basic data types

Name	Size (bits)	Usage	Range
bit	1	unsigned	0...1
bool	1	unsigned	0...1
byte	8	unsigned	0...255
short	16	signed	$-2^{15} - 1 \dots 2^{15} - 1$
int	32	signed	$-2^{31} - 1 \dots 2^{31} - 1$

Variables and Basic Data Types.....

- Like all modern well-structured programming languages, Promela also requires
 - that variables must be declared before they can be used
 - variable declarations follow the style of the C programming language ,
 - i.e. a basic data type followed by one or more identifiers and optional initializer
- byte count, total = 0;
- by default all variables of the basic types are initialized to 0.
 - as in C, 0 (zero) is interpreted as false while any non-zero value is interpreted as true .

The init process

- All Promela programs must contain an **init** process
 - is similar to the `main()` function within a C program
 - i.e. the execution of a Promela program **begins with the init process.**
- An **init** process
 - takes the form :

```
init { /*local declarations and statements*/ }
```
 - what could it do ???/
 - `init { skip }`
- While a **proctype** definition declares the behaviour of a process,
 - the instantiation and execution of a process definition is co-ordinated via the init process.

Statements

- A process body consists of sequence of statements.
- A statement is either
 - executable : it can be executed immediately
 - blocked : it cannot be executed
- An assignment statement is always executable

x=2 ;

Statements (contd)

- An expression is also a statement;

- it is executable if it evaluates to non-zero

skip, $2 < 3$ always executable

$x < 27$ executable if the value of x is less than 27

- A run statement is executable

- only if the process can be created
 - returns zero if this cannot be done
 - value otherwise returned is a run-time process ID number
 - run() is defined as an operator and so can be embedded in other expressions.

Illustrations

- run operator can pass parameter values to a new process

```
proctype A(byte state; short set)
{   (state==1) → state = set
}
init {run A(1,3)}
```

- run allows more than one copies of a process to be created

```
byte state = 1;
proctype A() { (state==1) → state = state + 1}
proctype B() { (state==1) → state = state - 1}
init { run A(); run B() }
```

Executability of Statements

■ Promela

- does not make a distinction between a condition and a statement
 - e.g. the simple boolean condition $a == b$ represents a statement in Promela.
- statements are either executable or blocked.
 - the execution of a statement is conditional on it not being blocked.
- notion of statement executability provides the basic means by which process synchronization can be achieved.

■ e.g.

```
while (a != b) skip /* conventional busy wait */  
(a == b) /* Promela equivalent */
```

Hello World

- A simple two process system

```
proctype hello() { printf("Hello") }
proctype world() { printf("World\n") }
init { run hello(); run world () }
```

- What does init do here ? What does run do here ?

- The run operator is executable only
 - if process instantiation is possible.
 - If a run is executable then a pid is returned.
 - the pid for a process can be accessed via the predefined local variable _pid.

- The execution of run does not wait

- for the associated process to terminate
 - i.e. further applications of run will be executed concurrently.

Process instantiation

- A process can be instantiated also by using active in front of proctype definition i.e.
 - HelloWorld can also be instantiated as

```
active proctype hello() { printf("Hello") }
active proctype world() { printf("World\n") }
```
- Multiple instances of the same proctype declaration can be generated using an optional array suffix, e.g.

```
active [4] proctype hello() { printf("Hello") }
active [7] proctype world() { printf("World\n") }
```

 - will generate 4 instances of hello and 7 instances of world.

PROMELA : An illustration...

```
proctype Foo(byte x) {  
    ...  
}  
  
init {  
    int pid2 = run Foo(2);  
    run Foo(27);  
}  
  
active [3] proctype Bar() {  
    ...  
}
```

- a process can be created at any point in the execution (even within any process)
- processes can also be created using active in front of proctype declaration

PROMELA – an illustration....

```
byte state = 2;  
  
proctype A() {  
    (state == 1) → (state = 3)  
}  
  
proctype B() {  
    state = state - 1  
}
```

; & → are statement separators

No init process
OR
“active” in
front of proctype
definition

Other Data Types

■ Arrays

- an array type is declared as `int table[max]`
- this generates an array of `max-1` integers
 - i.e. `table[0], table[1], ... table[max-1]`

■ Enumerated Types

- a set of symbolic constants is declared as
- `mtype = {LINE_CLEAR, TRAIN_ON_LINE, LINE_BLOCKED}`
 - a program can only contain one mtype declaration which must be global.

■ Structures

- a record data type is declared as `typedef`
`msg {byte data[4], byte checksum}`
- Structure access is as in C
 - `msg message;`
 - `... message.data[0]`

PROMELA – An illustration...

```
/* a Hello World PROMELA model for SPIN */
active proctype Hello() {
    printf("Hello process, my PID is : %d", _pid) ;
}

init {
    int lastPID;
    printf("init process, my PID is : %d", _pid);
    lastPID = run Hello();
    printf("last PID was : %d", _pid);
}
```

- How many processes are created, here ?

PROMELA – An illustration...Sample output

random seed

```
$ spin -n2 hello.pr _____  
init process, my pid is: 1  
last pid was: 2  
Hello process, my pid is: 0  
Hello process, my pid is: 2  
3 processes created
```

running SPIN in
random simulation mode

Statement Delimiters

- Two types of statement delimiters
 - ; and ->
 - use the one that is most appropriate at the given situation
 - usually, ; is used between ordinary statements
 - -> is often used after "guards" in a if OR do statement, pointing at what comes next
 - these can be used interchangably:

Statements: Reviewing

- The **skip** statement is always executable.
 - “does nothing”, only changes process’ process counter
- A **run** statement is only executable if a new process can be created (remember: the number of processes is bounded).
- A **printf** statement is always executable (but is not evaluated during verification, of course).

```
int x;
proctype Aap()
{
    int y=1;
    skip;
    run Noot();
    x=2;
    x>2 && y==1;
    skip;
}
```

Executable if **Noot** can
be created...

Can only become executable
if a **some other process**
makes **x** greater than 2.



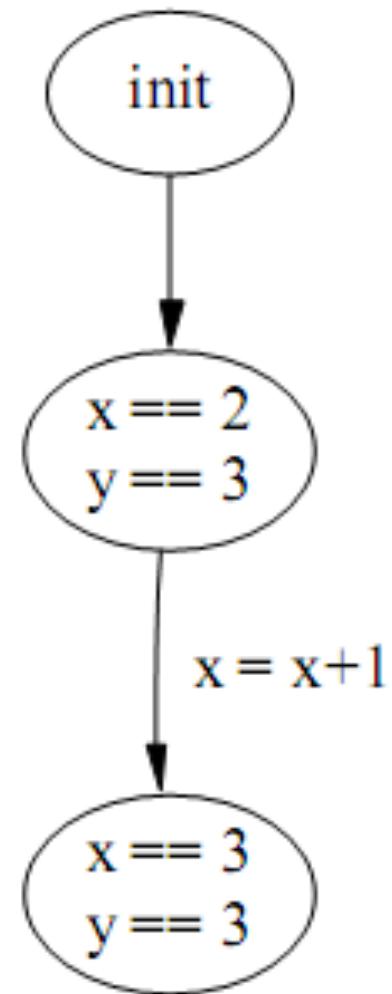
Interleaving Semantics

- Promela processes execute concurrently.
- Non-deterministic scheduling of the processes.
- Processes are interleaved
 - statements of different processes do not occur at the same time.
 - exception: rendezvous communication.
- All statements are atomic;
 - each statement is executed without interleaving with other processes.
- Each process may have several different possible actions enabled at each point of execution
 - only one choice is made, non-deterministically i.e. randomly

Processes as Automata

```
byte x = 2, y = 3;  
proctype A() {x = x + 1}  
proctype B() {  
    x = x - 1;  
    y = y + x  
}
```

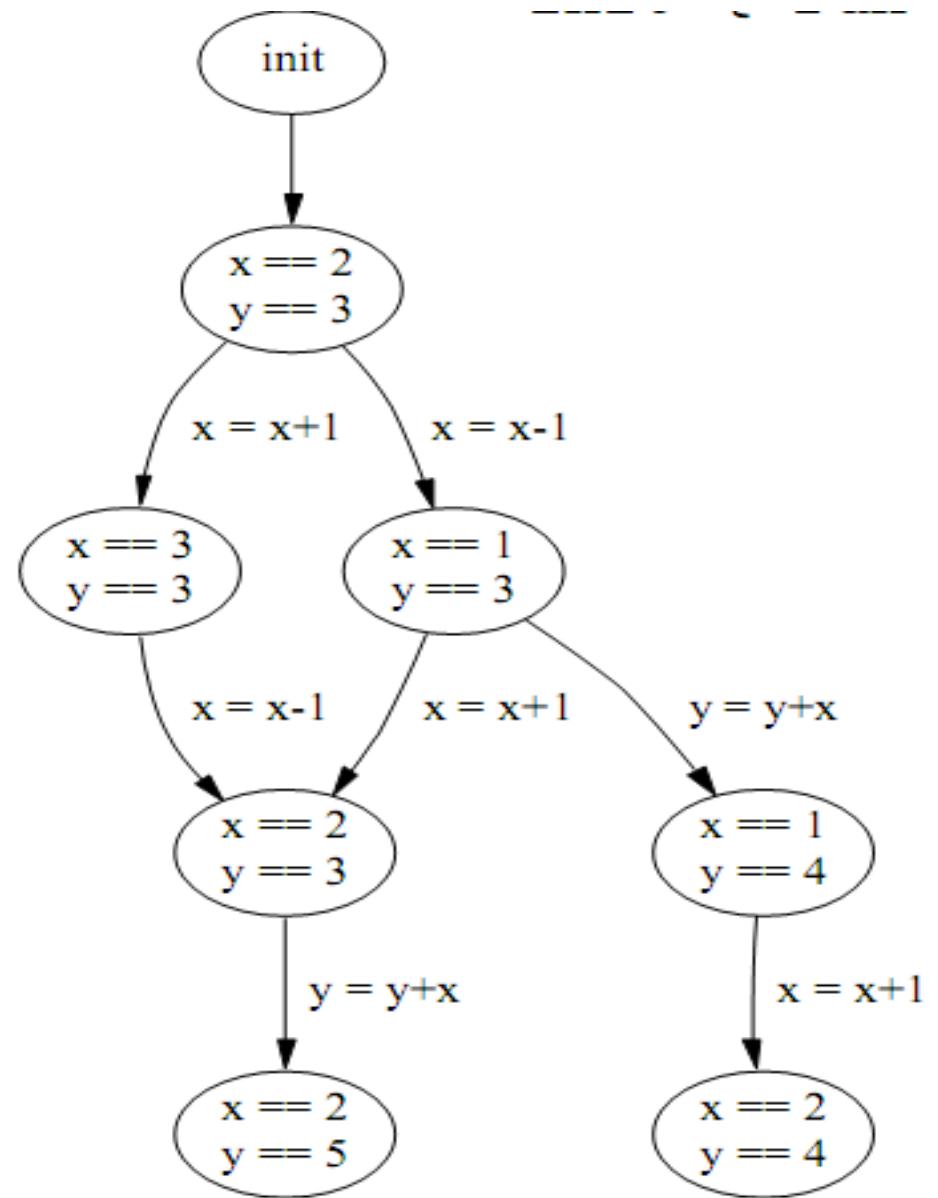
init {run A}; would cause this



Interleaved Execution

```
init {run A(); run B();}
```

```
byte x = 2, y = 3;  
proctype A(){x = x + 1}  
proctype B(){  
    x = x - 1;  
    y = y + x  
}
```



Deterministic vrs Nondeterministic behaviour

■ Deterministic behaviour

- a process is deterministic if for a given start state.....it behaves in exactly the same wayif supplied with the same stimuli from its environment.

■ Non-deterministic behaviour

- a process is non-deterministic if it need not always behave in exactly the same way..... each time it executes from a given start state withthe same stimuli from its environment.

■ Hence, race conditions can occur.....

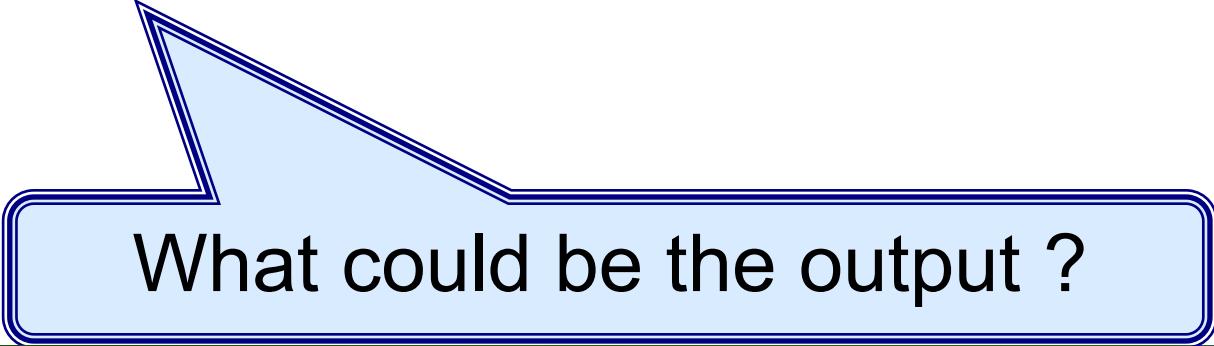
■ What are race conditions....?

Race conditions

■ Solutions

- use standard mutex algorithms
- use atomic sequences

```
byte state = 1;  
  
proctype A() {           { (state==1) → state=state + 1 } }  
proctype B() {           { (state==1) → state=state - 1 } }  
  
init { run A(); run B() }
```



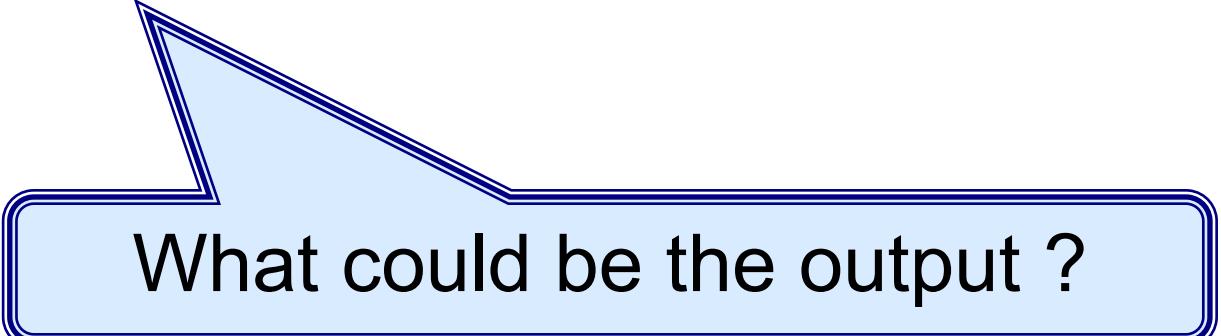
What could be the output ?

Race conditions

■ Solutions

- use standard mutex algorithms
- use atomic sequences

```
byte state = 1;  
  
proctype A() {atomic{ (state==1) → state=state + 1 } }  
proctype B() {atomic{ (state==1) → state=state - 1 } }  
init { run A(); run B() }
```



What could be the output ?

Another illustration: Atomic statements

```
proctype nr(short pid, a, b)
{ int res;
atomic { res = (a * a + b)/2*a;
         printf("result %d: %d\n", pid,
                res)
}
init { run nr(1,1,1); run nr(1,2,2); run
nr(1,3,2) }
```

Using atomic statements...

```
byte value1 = 1, value2 = 2, value3 = 3;

proctype A() { value3 = value3 + value2;
               assert( value3 == 5 )

}

proctype B() { value2 = value2 + value1;
               assert( value3 == 5 )

}

init { atomic{ run A(); run B() } }
```

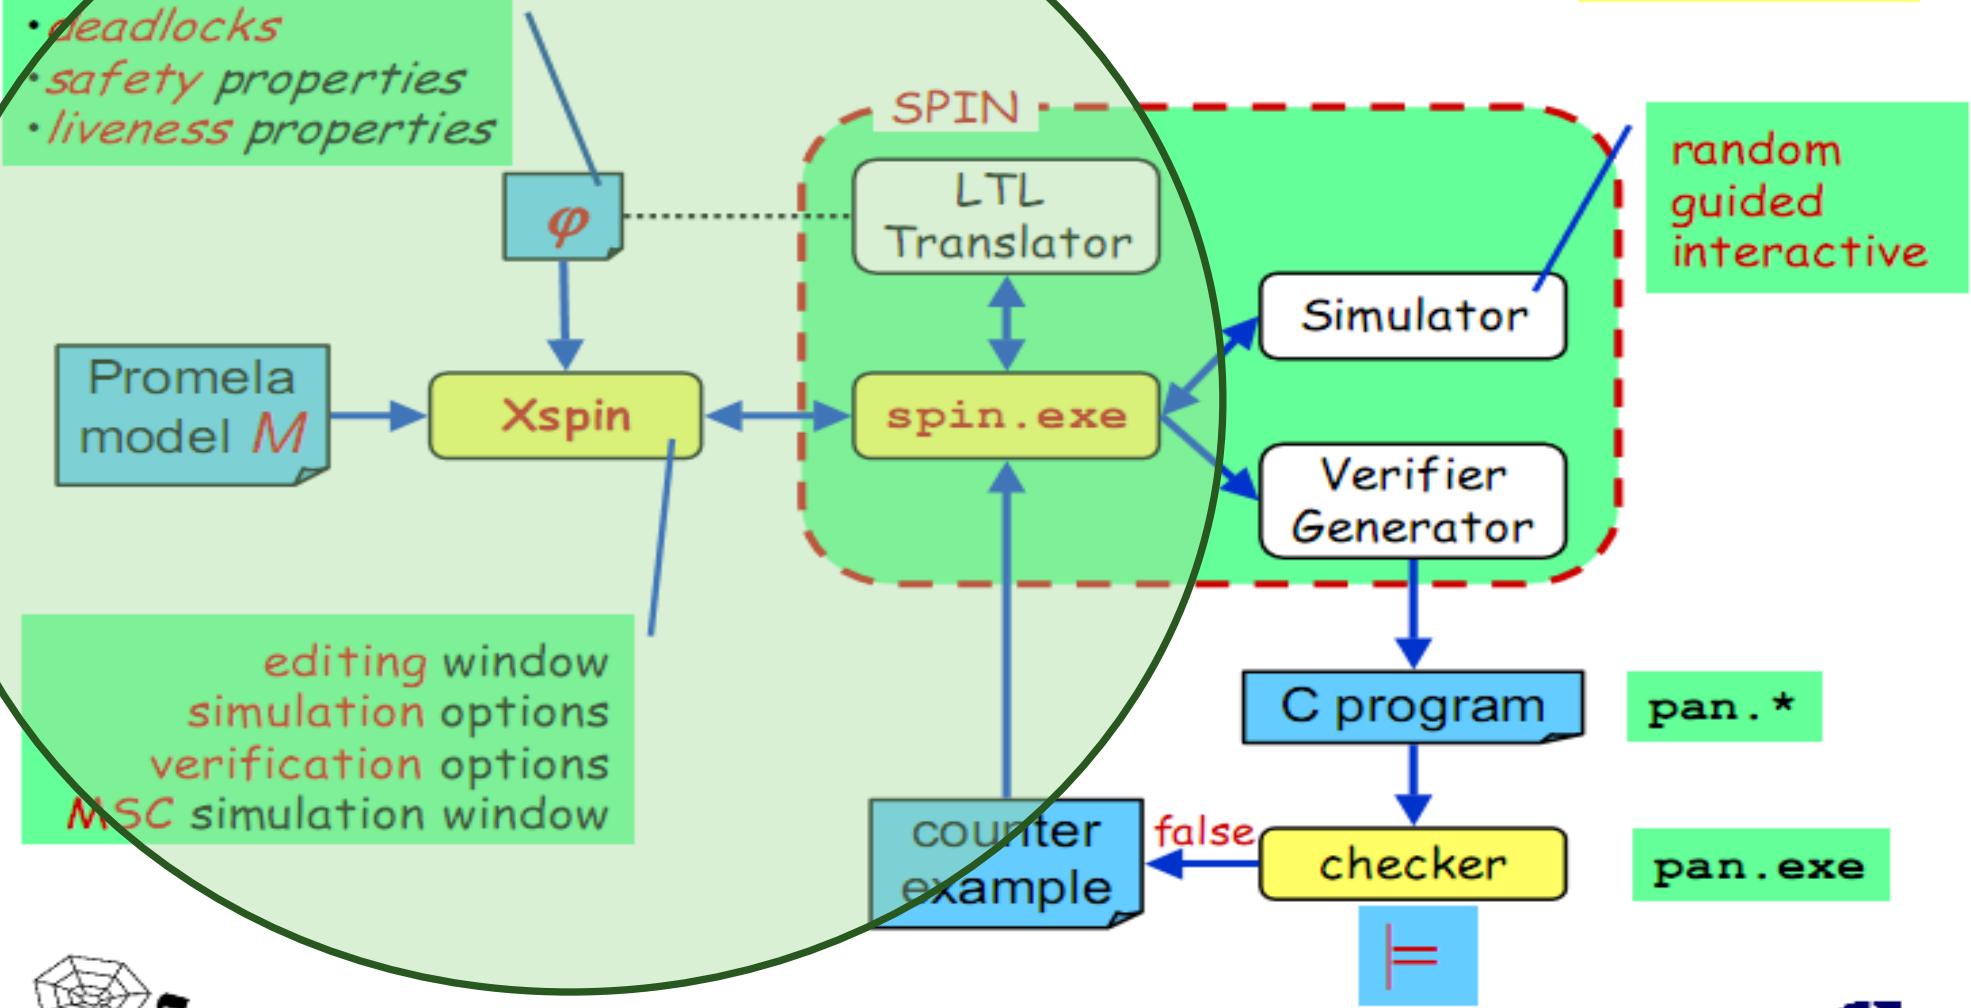
- When can **value3** here not be 5 – i.e. the assertion is violated ?

Using Atomic

- **atomic** keyword
 - helps avoid the undesirable interleaving of the PROMELA execution sequences....
 - restricts the level of interleaving and so
 - reduces complexity when it comes to validating a PROMELA model.
- However, **atomic** should be used carefully....Why ?

The XSPIN architecture

(X)SPIN Architecture



PROMELA Control Structures

- We have studied three ways for achieving control flow
 - Statement sequencing
 - Atomic sequencing
 - Concurrent process execution.
- Promela supports three additional control flow constructs
 - Case selection
 - Repetition
 - Unconditional jumps

Case Selection

```
byte count;  
proctype counter()  
{  
    if  
        :: count = count + 1  
        :: count = count - 1  
    fi  
}
```

- chooses one of the executable choices.
- If no choice is executable, the if-statement is blocked.
 - The **executability of the first statement (guard)** in each sequence determines **whether sequence is executed OR not**

Case Selection: An illustration

- An example of case selection with guards

if

```
:: (n % 2 != 0) -> n = n + 1;
```

```
:: (n % 2 == 0) -> skip;
```

fi

- If there is at least one choice (guard) executable,
 - the if statement is executable and SPIN non-deterministically chooses one of the alternatives.
- The operator “->” is equivalent to “;”.
 - By convention, it is used within if-statements to separate the guards from the statements that follow the guards.

Case Selection: An illustration....

- Guards need not be mutually exclusive

if

```
::: (x >= y) -> max = x;
```

```
::: (y >= x) -> max = y;
```

fi

- If **x** and **y** are equal then

- the selection of which statement sequence is executed is decided at random, giving rise to non-deterministic choice.

Repetition

- An example of repetition involving two statement sequences

do

```
::: (x >= y) -> x = x - y; q = q + 1;
```

```
::: (y > x) -> break;
```

od

- do statement is similar to the if statement.....

- however, instead of executing a choice ones, it keeps repeating the execution.
 - the (always executable) break statement may be used to exit a do-loop statement and transfers control to the end of the loop.

Repetition

- The first statement sequence denotes the body of the loop
 - while the second denotes the termination condition.
 - termination, however, is not always a desirable property of a system , in particular, when dealing with **reactive systems** :

do

```
:: (level > max) -> outlet = open;
```

```
:: (level < min) -> outlet = close;
```

od

What is a
reactive system ?

Repetition : An illustration

```
byte count;  
proctype counter()  
{
```

```
do  
:: count = count + 1  
:: count = count - 1  
:: (count == 0) -> break  
od  
}
```

Name	Size (bits)	Usage	Range
bit	1	unsigned	0...1
bool	1	unsigned	0...1
byte	8	unsigned	0...255
short	16	signed	$-2^{15} - 1 \dots 2^{15} - 1$
int	32	signed	$-2^{31} - 1 \dots 2^{31} - 1$

Combining Guards, Case, Repetition for purpose



```
proctype counter()
{
    do
        :: (count != 0) →
            if
                :: count = count + 1
                :: count = count - 1
            fi
        :: (count == 0) → break
    od
}
```

Unconditional Jump

- Promela supports the notion of an unconditional jump via the goto statement.

do

```
  :: (x >= y) -> x = x - y; q = q + 1;  
  :: (y > x) -> goto done;
```

od;

done:

skip

- “done” denotes a label.
 - a label can only appear before a statement.
 - a goto, like a skip, is always executable.

Assertions

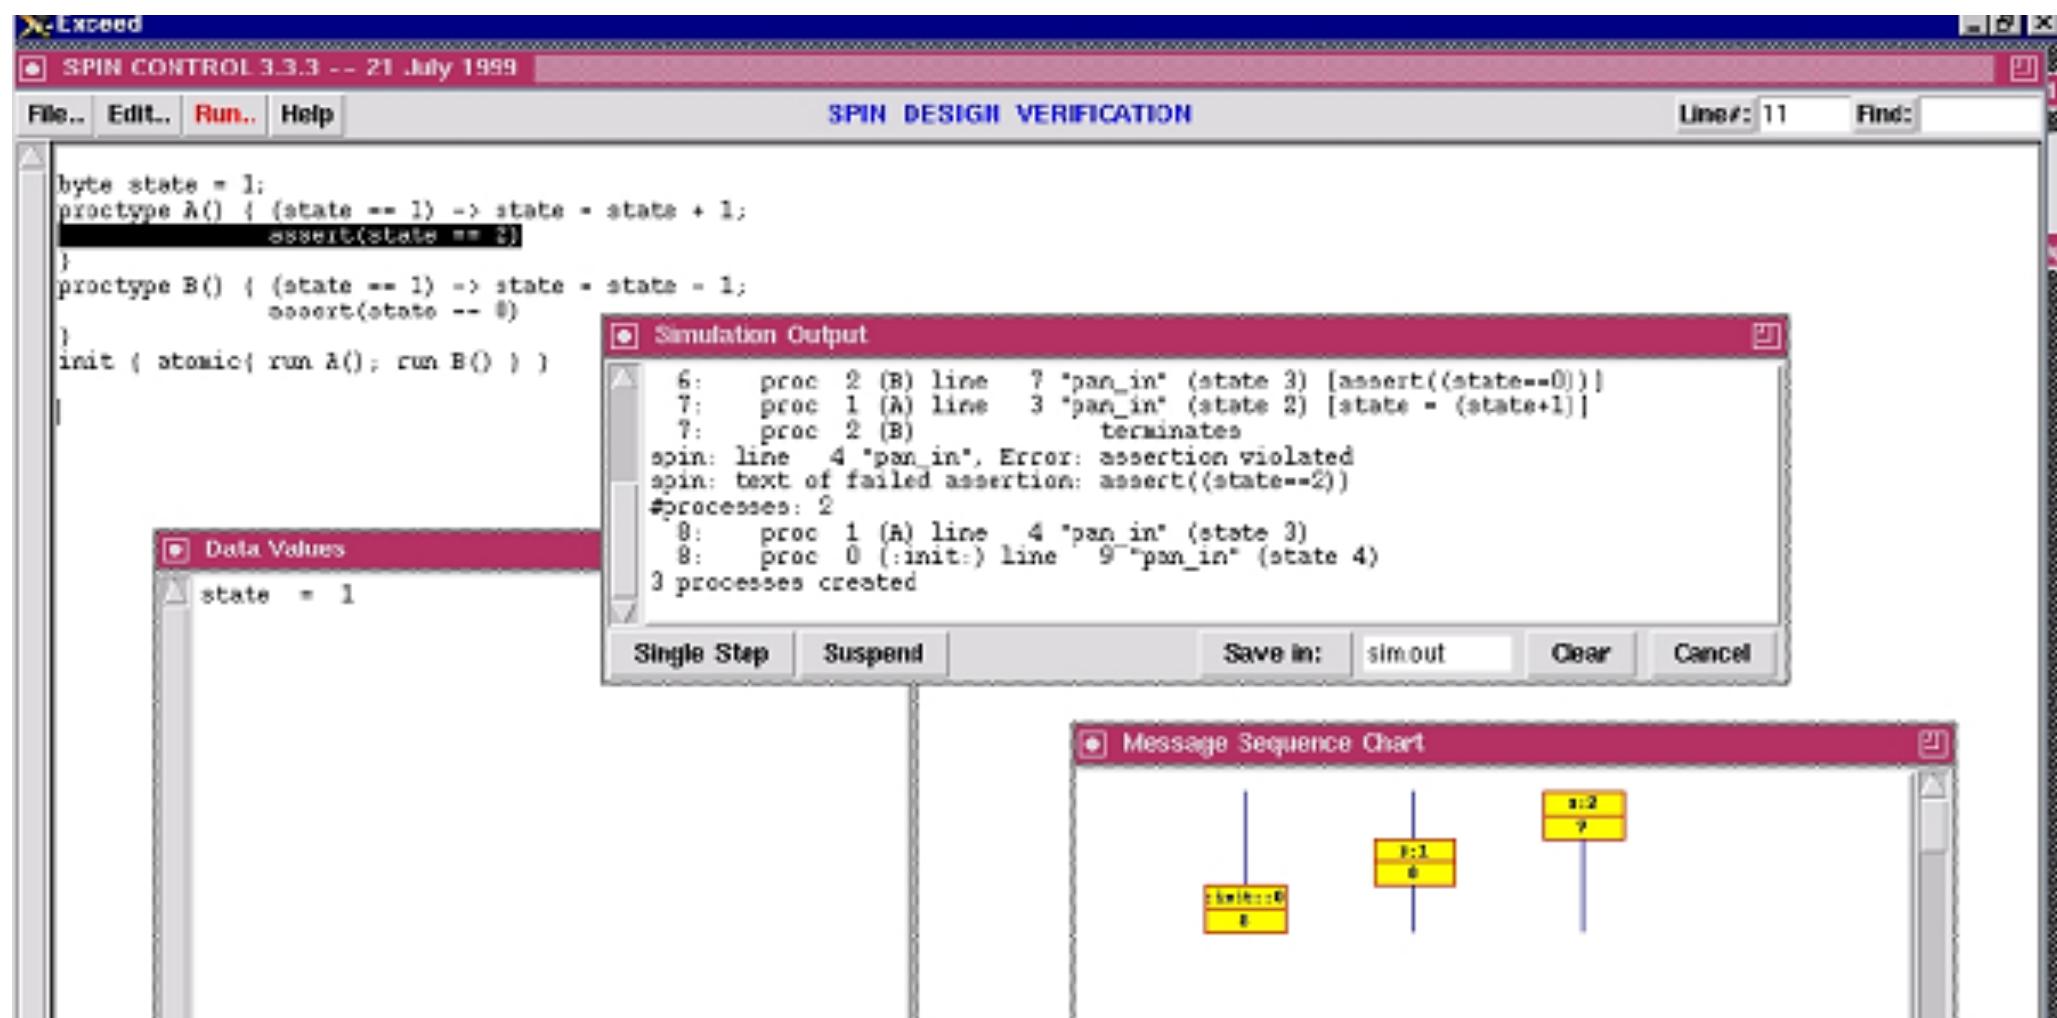
- An assertion is a statement which can be either true or false.
- Interleaving assertion evaluation with code execution provides
 - a simple yet very useful mechanism for checking desirable as well as erroneous behaviour with respect to our models.
- Assertion: Syntax within Promela
 - `assert(<logical-statement>)`
 - e.g. `assert(! (doors == open && lift == moving))`
- Within Promela we can express local assertions as well global system assertions.

Assertions: Example of Local Assertion

```
byte state = 1;  
  
proctype A() {  
    (state == 1) -> state = state + 1;  
    assert(state == 2)  
}  
  
proctype B() {  
    (state == 1) -> state = state - 1;  
    assert(state == 0)  
}  
  
init { atomic{ run A(); run B() } }
```

Will the assertion checking succeed here or fail?

Assertions...



Assertions Global Assertions

- A global assertion is also known as a system invariant
 - is a property that is true in the initial system state and remains true in all possible execution paths.
- To express a system invariant within Promela
 - one must define a monitor process that contains the desired system invariant.
- To ensure that the global assertion is checked anypoint during the execution....
 - an instance of the monitor process has to be run along with the rest of the system model
- In the case of a simulation the checking is not exhaustive, this is achieved within verification mode.

Mutual Exclusion#1

```
models: while (flag == 1) /*wait*/  
bit flag;  
byte mutex;  
proctype P(bit i) {  
    flag != 1;  
    flag = 1;  
    mutex++;  
    printf("MSC: P(%d) has entered the critical section\n, i);  
    mutex--;  
    flag = 0;  
}  
proctype monitor() {  
    assert(mutex != 2);  
}  
init { run P(0); run P(1); run monitor(); } }
```

- What could be the eventual value of mutex ?
- Is it really achieved ?
- Is the assertion preserved or violated, here ?

Mutual Exclusion#1

```
bit flag;
byte mutex;
proctype P(bit i) {
    flag != 1;
    flag = 1;
    mutex++;
    printf("MSC: P(%d) has entered the critical section\n", i);
    mutex--;
    flag = 0;
}
proctype monitor() {
    assert(mutex != 2);
}
init { atomic {run P(0); run P(1); run monitor(); } }
```

- What could be the eventual value of mutex ?
- Is it really achieved ?
- Is the assertion preserved or violated, here ?

Mutual Exclusion#2

```
bit flag1, flag2;  
byte mutex;
```

```
active proctype A() {  
    flag1 = 1;  
    flag2 == 0;  
    mutex++;  
    mutex--;  
    flag1 = 0;  
}
```

```
active proctype B() {  
    flag2 = 1;  
    flag1 == 0;  
    mutex++;  
    mutex--;  
    flag2 = 0;  
}
```

```
active proctype monitor() { assert mutex != 2);
```

- What could be the eventual value of mutex ?
- Is it really achieved ?
- Is the assertion preserved or violated, here ?

“Invalid End state” in SPIN

Mutual Exclusion#3 - Dekker's algorithm

```
bit flag1, flag2; byte mutex, turn;
```

```
active proctype A() {
    flag1 = 1;
    turn = B_TURN;
    flag2 == 0 || turn == A_TURN;
    mutex++;
    mutex--;
    flag1 = 0;
}
```

```
active proctype B() {
    flag2 = 1;
    turn = A_TURN;
    flag1 == 0 || turn == B_TURN;
    mutex++;
    mutex--;
    flag2 = 0;
}
```

```
active proctype monitor() { assert mutex != 2);
```

First software-only solution to the mutex problem for two processes...

Dekker's algorithm

```
#define true 1 false 0 Aturn 1 Bturn 0
bool ARuns, BRuns, t;

proctype A()
{
    ARuns = true;
    t = Bturn;
    (BRUNS==false || t==Aturn);
    /*critical section */
    ARUNS = false
}

proctype B()
{
    BRUNS = true;
    t = Aturn;
    (ARUNS==false || t==Bturn);
    /*critical section */
    BRUNS = false
}

init { run A(); run B(); }
```

Can this be generalized to a single process?

Identify the output

```
int a, b, quo, rem;
bit done, load;

proctype quo_rem()
{
    do
        :: (load == 1) -> quo = 0; rem = a; done = 0;
        :: (load != 1) -> if
            :: (rem >= b) -> rem = rem-b; quo = quo+1
            :: (b > rem)   -> done = 1
            fi
    od
}

init { a = 7; b = 2; done = 1; load = 1;
       run quo_rem();
       done == 0; load = 0; done == 1;
       printf("result1 = %d\n", quo);
       printf("result2 = %d\n", rem)
}
```

Identify the difference and the output

```
byte in1, in2, a, b, quo, rem;
bit load = 0, done = 1;

proctype quo_rem() {
    do
        :: (load == 1) -> a = in1; b = in2;
                    quo = 0; rem = a; done = 0;
        :: (load != 1) -> if
            :: (rem >= b) -> rem = rem-b; quo = quo+1
            :: (b > rem) -> done = 1
        fi
    od
}

proctype env() {
    in1 = 7; in2 = 2; load = 1;      /* init inputs */
    done == 0; load = 0; done == 1;  /* read results */
    printf("quotient = %d\n", quo);
    printf("remainder = %d\n", rem) }

init { atomic{ run quo_rem(); run env() } }
```

Timeouts

- Reactive systems typically require a means of aborting OR rebooting when a system deadlocks.
- Promela provides a primitive statement called timeout for the purpose e.g.

```
proctype watchdog ()  
{  do  
    :: timeout -> guard!reset  
  od  
}
```

- The timeout condition **becomes true** when **no other statements** within the overall system being modelled are **executable**.

Exceptions

- the **unless** statement

- a useful exception handling feature

```
{ statements-1 }
```

```
unless
```

```
{ statements-2 }
```

- Consider an alternative watchdog process

```
proctype watchdog ()
```

```
{
```

```
do
```

```
:: process_data() unless guard?reset;
```

```
:: process_reset()
```

```
od
```

```
}
```

Message Channels

- What means have been studied so far to achieve communication between distinct processes ?
- However, Promela supports **message channels** also
 - provide a more natural and sophisticated means of modeling inter-process communication / data transfer.

Message Channels....

- a channel declaration could take any of the forms....

```
chan <name> =<dim> of {<t1>,<t2>,<t3> ...<tn>}
```

- Semantics

```
channel = FIFO-buffer (for dim>0)
```

- e.g.

```
chan q = [1] of {byte}
```

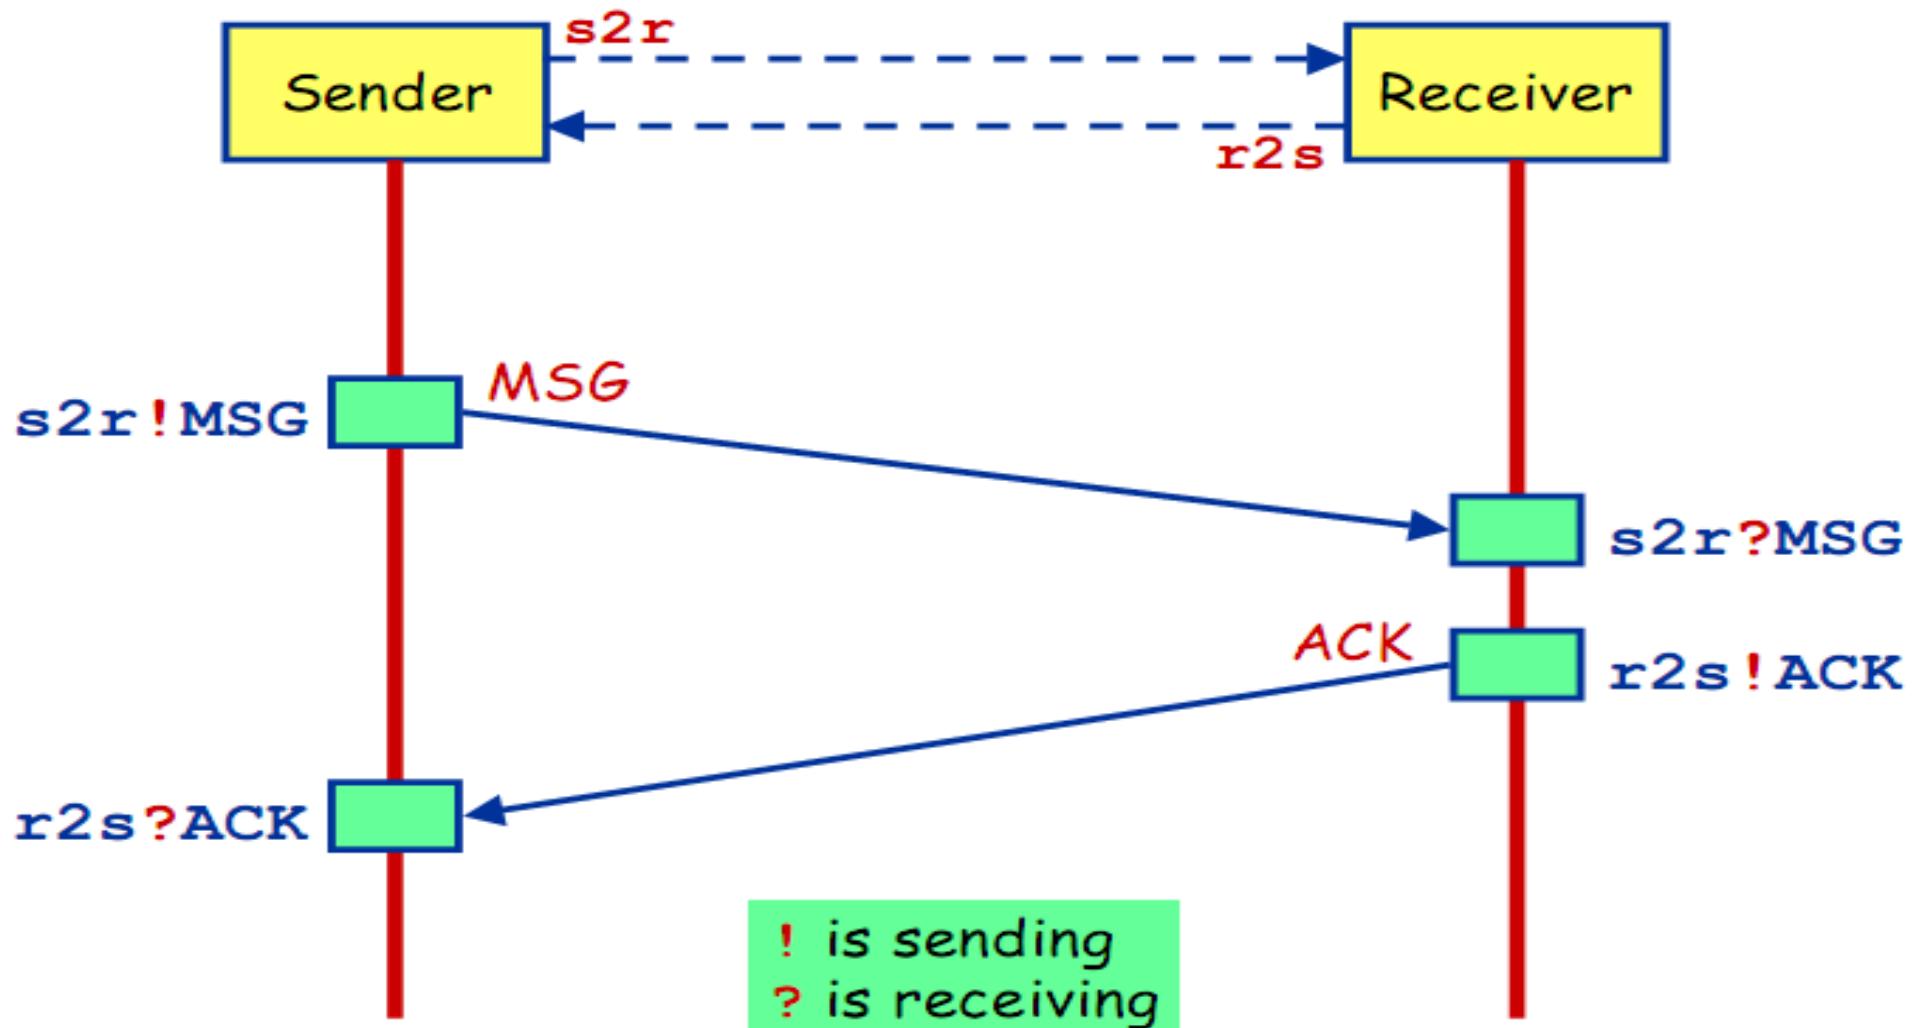
```
chan q = [3] of {mtype, int}
```

- if dim=0 then synchronous

```
chan q = [0] of {bit}
```

- A channel can be defined to be either local or global.

Message Channels....



Message Channels....Sending-in

- Sending messages through a channel - FIFO buffer (`dim >0`)
 - achieved by the `!` operator,
 - e.g. `in_data ! 4;`
 - type of the channel and variable must match
- For multiple data values are to be transferred via each message
 - `out_data ! x + 1, true, in_data;`
where `x` is of type byte.
- The executability of a send statement is dependent upon the associated **channel being non-full**.

Message Channels....Receiving from

- Receiving messages is achieved by the ? operator
e.g. `in_data ? msg;`
- if the channel is not empty, the first message is fetched from the channel and the is stored `in_data`
- Multiple values can also be fetched
e.g. `out_data ? value1, value2, value3;`
- the executability of a receive statement is dependent up on the associated channel being non-empty,
 - e.g. the following statement will be blocked
`in_data ? value;`
unless `in_data` contains at least one message .

An illustration : Message Channels...

```
#define N 128
#define SIZE 16

chan in = [size] of {short};
chan large = [size] of {short};
chan small = [size] of {short};

active proctype split() {
    short cargo;
    do
        ::in ? cargo →
            if
                :: (cargo >= N) → large!cargo
                :: (cargo < N) → small!cargo
            fi
    od
}
```

Message Channels....

- If more data values are sent per message than can be stored by a channel then the extra data values are lost

e.g. `in_data ! msg1, msg2;`

the msg2 will be lost.

- If fewer data values are sent per message than are expected then the missing data values are undefined,

e.g. `out_data ! 4, true;`

`out_data ? x, y, z;`

x and y will be assigned the values 4 and true respectively while the value of z will be undefined.

Message Channels....

■ **len** operator

- to determine the number of messages in a channel
 - e.g. `len(in_data)`
- if the channel is empty then the statement will block.

■ **empty**, **full** operators

- determine whether or not messages can be received or sent respectively,
 - e.g. `empty(in_data); full(in_data)`

■ Non-destructive retrieve

`out_data ? [x, y, z]`

returns 1 if `out_data?x,y,z` is executable otherwise 0.

- No side-effects – evaluation, not execution, i.e. no message retrieved.

Channels as parameters

```
proctype A(chan q1) {  
    chan q2;  
    q1?q2;  
    q2!123  
}  
  
proctype B(chan qforb) {  
    int x;  
    qforb?x;  
    printf("x = %d\n", x)  
}  
  
init {  
    chan qname[2] = [1] of { chan } ;  
    chan qforb = [1] of { int } ;  
    run A(qname[0]); run B(qforb);  
    qname[0]!qforb  
}
```

What will be
the output ?

Channels as parameters....

```
proctype A(chan q1) {
    chan q2;
    q1?q2;
    q2!99
}
proctype B(chan qforb) {
    int x;
    qforb?x;
    x++;
    printf("x == %d\n", x)
}
init {
    chan qname = [1] of { chan };
    chan qforb = [1] of { int };
    run A(qname); run B(qforb); qname!qforb
}
```

What will be
the side-effect
of running
this program?

Channels as parameters :Division

```
proctype division(int x,y,q; chan res) {
    if
        :: (y > x) -> res!q,x;
        :: (x >= y) -> run division(x - y, y, q + 1, res);
    fi
}
init{
    int q,r;
    chan child = [1] of { int, int };
    run division(7, 3, 0, child);
    child ? q,r;
    printf("result: %d %d\n", q,r)
}
```

Note that
this is a
tail-recursive
program

process factorial

```
proctype fact(int n, chan p)
{ int result;
  if
    :: (n <=1) → p ! 1
    :: (n >=2 →
        chan child = [1] of {int};
        run fact(n-1, child);
        child ? result;
        p ! n *result
      fi
}
```

Reason why is it
NOT a
tail-recursive
program

```
init {
  int result;
  chan child = [1] of {int};
  run fact(7,child);
  child ? result;
  printf("result:%d\n", result)
}
```

Communication type

- What was the type of communication pattern observed in the examples till now ? i.e.
 - Synchronous OR Asynchronous communication ?
 - Why ?

Synchronous communication

- when the channel declaration is
`chan ch = [0] of {bit, byte};` i.e. when `dim=0`
- if `ch!x` is enabled and
 - if there is a corresponding receive `ch?x`
 - that can be executed simultaneously and both the statements are enabled
 - both statements will handshake and together do the transition

`chan ch =[0] of {bit,byte};`

P wants to do `ch!1,3+7`

Q wants to do `ch?1,x`

then after communication x will be 10

Synchronous communication: An illustration

```
#define msgtype 33
chan name = [0] of { byte, byte }
proctype A() {
    name!msgtype(124);
    name!msgtype(121)
}
proctype B() {
    byte state;
    name?msgtype(state)
}
init { atomic { run A(); run B() } }
```

Rendezvous communication

```
#define msgtype 13
byte name;
chan name = [0] of {byte, byte}
proctype A()
{
    name ! msgtype(124);
    name ! msgtype(121)
}```

proctype B()
{
    byte state;
    name ? msgtype(state)
}

init
{ atomic {run A(); run B()} }
```

What are the implications of atomic?
What if the channel sizes were two? one?

Dijkstra's semaphores

```
#define p 0 v 1
chan sema = [0] of {bit};
proctype dijkstra () {
do
:: sema ! p → sema ? v
od
}
proctype user() {
    sema ? p;
        /*critical section */
    sema!v
}
init{
    atomic { run dijkstra(); run user(); run user();
run user(); }
}
```

Semaphores Revisited : What is the difference?

```
#define p 0 #define v 1
chan sema = [0] of { bit };

proctype semaphore() {
    do :: sema!p -> sema?v od
}

byte count;
proctype user() {
    sema?p;
    count = count + 1;
    skip; /* critical section */
    count = count - 1;
    sema!v; skip /* non-critical section */ }

proctype monitor() {
    do :: assert(count == 0 || count == 1) od
}

init { atomic {run monitor(); run semaphore(); run user(); run
user(); run user()} }
```

What does this protocol do ?

```
proctype Sender(chan in, chan out)
{
    bit sendbit, recvbit;
    do
        :: out ! MSG, sendbit ->
            in ? ACK, recvbit ;
            if
                :: recvbit == sendbit -> sendbit = 1-
                                            sendbit;
                :: else -> skip
            fi
    od
}
```

Alternating Bit Protocol

- To every message, the sender adds a bit.
- The receiver acknowledges each message by sending the received bit back.
- The receiver only accepts messages with a bit that it expected to receive.
- If the sender is sure that the receiver has correctly received the previous message, it sends a new message and it alternates the accompanying bit.

Alternating Bit protocol

```
mtype = {msg, ack};  
chan to_sndr = [2] of {mtype, bit};  
cha to_rcvr = [2] of {mtype, bit};  
proctype Sender(chan in, chan out)  
{  
    bit sendval, recval;  
    do  
        :: out!msg(sendval) ->  
            in?ack(recval);  
            if  
                :: recval == sendval ->  
                    sendval = 1 - recval  
                :: else -> skip  
            fi  
    od  
}
```

```
proctype Receiver(chan in, chan  
                  out)  
{  
    bit recval;  
    do  
        :: in!msg(recval) ->  
            out?ack(recval);  
        :: timeout ->  
            out?ack(recval);  
    od  
}  
  
init {  
    run Sender(to_sndr, to_rcvr);  
    run Receiver(to_rcvr, to_sndr);  
}
```

Modeling Loss of Messages

Reviewing Interleaving Semantics...

- Transition has two components
 - Side-effect free condition and an atomic action
- A transition is executable if its condition holds, otherwise it is blocked
 - Assignments are always executable
 - Run statements are executable if new processes can be created
 - Conditions are executable if the truth value is true
 - Send statements are executable if channel not full (or ...)
 - Receive statements are executable if channel is nonempty and patterns match
 - Skip statements are always executable

Traffic Lights

```
mtype = (RED, YELLOW, GREEN);  
active proctype TrafficLight() {  
    byte state = GREEN;  
    do  
        :: (state==GREEN) -> state=YELLOW;  
        :: (state==YELLOW) -> state=RED;  
        :: (state==RED) -> state=GREEN;  
    od;  
}
```

An illustration

```
#define train 1

chan BlockSecAB = [2] of { bit };
chan BlockSecBC = [2] of { bit };

proctype SignalA(chan out_track)
{
    do
        :: out_track!1
    od
}

proctype SignalB(chan in_track,
                 out_track)
{
    do
        :: in_track?train;
        out_track!train
    od
}
```

```
proctype SignalC(chan in_track)
{
    do
        :: in_track?train;
    od
}

init { atomic{
    run SignalA(BlockSecAB);
    run SignalB(BlockSecAB,
                BlockSecBC);
    run SignalC(BlockSecBC)
}
`
```

An illustration modified

```
chan TunnelAB = [2] of { byte } ;
chan TunnelBC = [2] of { byte } ;
chan TunnelCD = [2] of { byte } ;
chan TunnelDA = [2] of { byte } ;

proctype Station(chan in_track, out_track)
{
    byte train;

    do
        :: in_track?train; out_track!train
    od
}
```

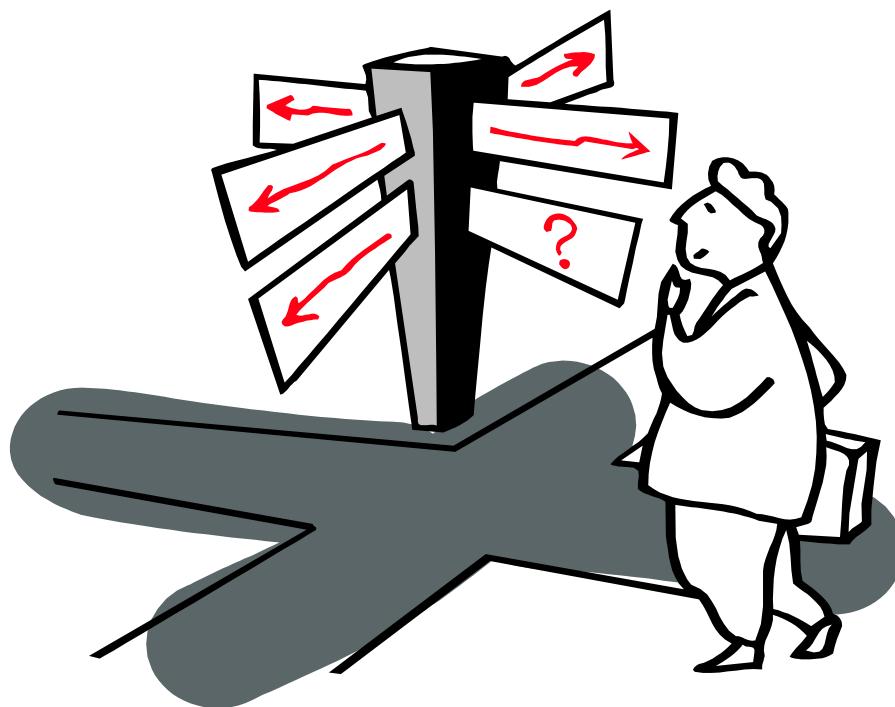
An illustration modified...

```
proctype Setup(chan track; byte train)
{
    track!train;
}

init { atomic{
    run Setup(TunnelBC, 1); /* introduce train 1 before station C */
    run Setup(TunnelDA, 2); /* introduce train 2 before station A */

    run Station(TunnelDA, TunnelAB); /* station A */
    run Station(TunnelAB, TunnelBC); /* station B */
    run Station(TunnelBC, TunnelCD); /* station C */
    run Station(TunnelCD, TunnelDA) } /* station D */
}
```


To teach is to learn twice !!



Hence, Thank You !!!